

Improving SOA Antipattern Detection in Service Based Systems by Mining Execution Traces

Mathieu Nayrolles, Naouel Moha and Petko Valtchev

LATECE Team, Département d'informatique, Université du Québec à Montréal, Canada
mathieu.nayrolles@gmail.com, {moha.naouel, valtchev.petko}@uqam.ca

Abstract—Service Based Systems (SBSs), like other software systems, evolve due to changes in both user requirements and execution contexts. Continuous evolution could easily deteriorate the design and reduce the Quality of Service (QoS) of SBSs and may result in poor design solutions, commonly known as SOA antipatterns. SOA antipatterns lead to a reduced maintainability and reusability of SBSs. It is therefore important to first detect and then remove them. However, techniques for SOA antipattern detection are still in their infancy, and there are hardly any tools for their automatic detection. In this paper, we propose a new and innovative approach for SOA antipattern detection called SOMAD (Service Oriented Mining for Antipattern Detection). SOMAD is an evolution of the only—and therefore—state-of-the-art approach named SODA. Our new approach goal is to improve SOA antipattern detection in SBSs by mining execution traces. SOMAD mines strong associations between sequences of service/method calls from the execution traces of a SBS and further filters them by means of a suite of dedicated metrics. We first discuss the underlying association mining model and the intuitions behind the SBS-oriented rule metrics. We then present a validating application of our approach to two independently developed SBSs. A comparison of our SOMAD implementation with SODA reveals superiority of the former: Its precision is better by a fairly comfortable margin ranging from 2.6% to 16.67% while keeping the recall to 100% and being—at least—2.5 times faster on the same test subjects.

Index Terms—SOA Antipatterns, Mining Execution Traces, Sequential Association Rules, Service Oriented Architecture.

I. INTRODUCTION

Service Based Systems (SBSs) are composed of ready-made services that are accessed through the Internet [1]. Services are autonomous, interoperable, and reusable software units that can be implemented using a wide range of technologies like Web Services, REST (REpresentational State Transfer), or SCA (Service Component Architecture, on the top of SOA.). Most of the world biggest computational platforms: Amazon, Paypal, and eBay, for example, represent large-scale SBSs. Such systems are complex—they may generate massive flows of communication between services—and highly dynamic: services appear, disappear or get modified. The constant evolution in an SBS can easily deteriorate the overall architecture of the system and thus make architectural defects, known as SOA *antipatterns* [2], appear. An antipattern is the opposite of a design pattern: while design patterns should be followed to create more maintainable and reusable systems [3], antipatterns should be avoided since they have a negative impact, e.g., hinder the maintenance and reusability of SBSs.

Given the negative impact of SOA antipatterns on main-

tainability and reusability, there is a clear and urgent need for techniques and tools to detect them. However, the highly dynamic and distributed nature of SBSs makes SOA antipatterns detection a challenge comparing to OO systems [4]–[6]. Recently, a tool was developed by our team, called SODA (Service Oriented Detection for Antipatterns) [2], [7], which targets the detection of SOA antipatterns. The tool relies on a Domain Specific Language (DSL) for specifying SOA antipatterns based on metrics (mostly static plus a few dynamic ones), on one hand, and on a method for automated generation of detection algorithms from antipattern specifications, on the other hand.

Albeit efficient and precise, SODA suffers from serious limitations. Indeed, SODA performs two phases of analysis, a first one, static, followed by a second one, dynamic. The first static analysis requires access to service interfaces. Consequently, SODA cannot analyze systems that are proprietary or not open-source. The second, and smaller, dynamic analysis requires the execution of the system and therefore, the creation of runnable scenarios. Moreover, SODA has been specifically designed for systems based on the SCA standard and its precision decreases as the target system gets bigger. Given SODA's limitations, there is a clear emergency to improve our tools and techniques in order to detect SOA antipatterns in other SOA technologies while improving our precision. In this article, we propose a new and innovative approach for the detection of SOA antipatterns named SOMAD (Service Oriented Mining for Antipattern Detection). This approach does not require scenarios as SODA to concretely invoke service interfaces and it only relies on execution traces that may be provided by any SOA technologies. SOMAD is able to eliminate the non-relevant data using data mining techniques—sequential association rules mining—and then, detect antipatterns. It discovers SOA antipatterns by first extracting associations between services as expressed in the execution traces of an SBS. To that end, we apply a specific variant of the association rule mining task based on sequences or episodes: In our case the sequences represent services or, alternatively, method calls. Further, we filter these generated association rules by means of a suite of dedicated metrics.

The main contributions of this paper are twofold: (i) a new approach named SOMAD for the detection of SOA antipatterns. This approach is based on association rules mining and applies on execution traces of any kind of SBSs from different SOA technologies; (ii) a empirical validation of this approach, which outperforms the only other tool for the detection of

SOA antipatterns in terms of precision (2.6% to 16.67%) and efficiency (2.5 times faster) and therefore, improving the detection of SOA antipatterns.

We applied SOMAD on two different SBS called *Home Automation* and *FraSCaTi* [8]. *Home Automation* is made of 13 services and *FraSCaTi* is almost ten times larger. We compared the outcome of SOMAD to the one produced by SODA, the so far unique tool for antipattern detection from the literature. Both tools were evaluated in terms of precision and recall, on one hand, and efficiency, on the other hand. The study results indicate that SOMAD significantly outperforms SODA in term of precision and efficiency.

The remainder of the article is organized as follows. Section II presents related works on pattern and antipattern detection both in SOA and OO paradigms and related works on knowledge extraction. Section III presents the SOMAD approach, and in particular the mining of association rules from execution traces while, Section IV presents our experimental study with a comparison of our SOMAD approach to SODA. Finally, we provide some concluding remarks in Section V.

II. RELATED WORK

Our approach stands on two different fields: antipattern detection and knowledge extraction from execution traces; this section is organized as follows. Section II-A presents related works on the pattern and antipattern detection both in OO and SOA paradigms while Section II-B presents related works on knowledge extraction. Finally, Section II-C presents the SODA-state-of-the-art approach we proposed in [2], [7] in order to make a fair comparison with SOMAD.

A. Pattern and antipattern detection

Architectural (or design) quality is essential for building well-designed, maintainable, and evolvable SBSs. Patterns and antipatterns have been recognized as one of the best ways to express architectural concerns and solutions, and thus target high quality in systems. A number of methods and tools exist for the detection of antipatterns in OO systems [5], [9], [10] whereas the relevant theory and practices have been summarized in recent books [11], [12]. However, unlike OO antipatterns, the detection of their SOA counterparts is still in its infancy. An alternative approach to the declarative specification of antipatterns, called SPARSE, is presented in [13]. In SPARSE, antipatterns are described as an OWL ontology augmented with a SWRL (Semantic Web Rule Language) rule basis whereas their occurrences are tested through automated reasoning.

Other relevant work has focused on the detection of specific antipatterns related to the system's performance and resource usage and/or given technologies. For example, Wong *et al.* [14] use a genetic algorithm for detecting software faults and anomalous behavior in the resource usage of a system (e.g. memory usage, processor usage, thread count). The approach is driven by *utility functions* that correspond to predicates identifying suspicious behavior by means of resource usage metrics. In another related work, Parsons *et al.* [15] tackled the detection of performance antipatterns. They use a rule-based approach made of both static and dynamic analyzes that are

tailored to component-based enterprise systems (in particular, JEE applications).

B. Knowledge extraction

A large number of studies focused on knowledge extraction from execution traces. They were motivated by the identification of: crosscutting concerns (aspects) [16], business processes [17], patterns of interests among service users [18], [19], and features either in OO systems [20] or SBSs [21]. Further related work focused on the identification of service composition patterns [22], i.e. sets of services that are repetitively used together in different systems and that are structurally and functionally similar. These patterns correspond to good practices in designing and developing SBSs. This work could be adapted to antipattern correction, e.g. of highly coupled *Tiny Services* to merge into a unique service.

Few projects have explored pattern detection through execution trace mining. Ka-Yee Ng *et al.* [23] proposed MoDeC, an approach for identifying behavioral and creational design patterns using dynamic analysis and constraint programming. They reverse-engineer scenario diagrams from an OO system by bytecode instrumentation and apply constraint programming to detect these patterns as runtime collaborations. Hu and Sartipi [24] tackle the detection of design patterns in traces using scenario execution, pattern mining, and concept analysis. The approach is guided by a set of feature-specific scenarios to identify patterns, as opposed to a general pattern detection.

Although different in goals and scope, the above studies on OO antipatterns form a sound basis of expertise and technical knowledge for building methods for the detection of SOA antipatterns. However, despite a large number of commonalities, OO (anti)pattern detection methods cannot directly apply to SOA. Indeed, SOA focuses on services as first-class entities and thus remains at a higher granularity level than OO classes. Moreover, the highly dynamic nature of a SBS raises challenges that are not preponderant in OO systems.

C. SODA : The state-of-the-art tool

As mentioned earlier, the only approach currently available is SODA [2], [7]. SODA relies on a rule-based language that enables antipattern specification using a set of metrics. A generic process then turns the specification into detection algorithms. Although efficient and precise, SODA is an intrusive approach because it requires a set of valid scenarios concretely invoking the interface methods of SBSs and its dynamic analysis involves SCA properties. SODA includes the following three main steps (as illustrated in Figure 1):

Specification of SOA antipatterns: This step lies in identifying properties relevant to the specification of SOA antipatterns. These properties mainly correspond to metrics including, for example, the cohesion, coupling, number of methods, response time and availability. Moreover, these properties are used as a base vocabulary to a DSL, in the form of a rule-based language for specifying SOA antipatterns through *rules cards*, which are sets of rules combining metrics.

Generation of detection algorithms: This step consists in generating automatically detection algorithms by visiting mod-

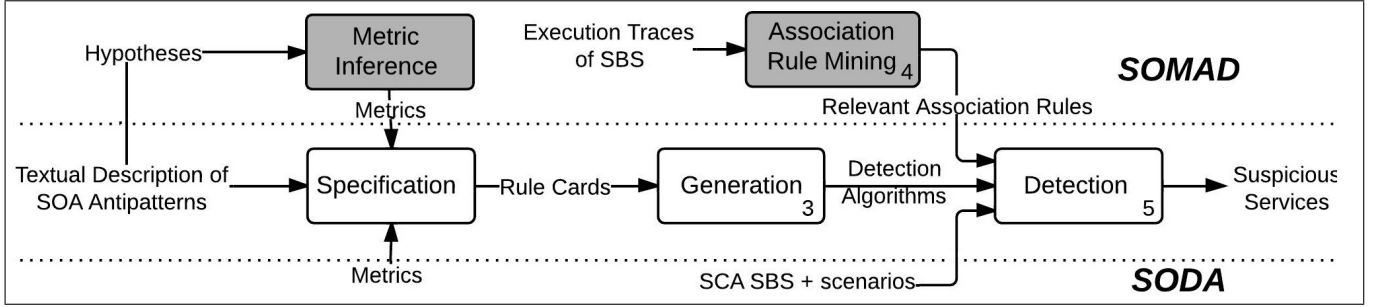


Figure 1: The SODA and SOMAD approaches. Grey boxes correspond to the new steps of SOMAD added to the SODA steps in white.

els of rule cards specified in the previous step. This process is straightforward and enables the automatic generation of detection algorithms directly executable.

Detection of SOA antipatterns: The third and last step of SODA consists in applying the detection algorithms generated in the previous step on the SBS of interest. This step allows the automatic detection of SOA antipatterns using a set of predefined scenarios to invoke service interfaces. At the end of this step, services in the SBS suspected to be involved in an antipattern are identified.

III. THE SOMAD APPROACH

We propose a five steps approach, named SOMAD (Service Oriented Mining for Antipattern Detection), for the detection of SOA antipatterns within execution traces of SBSs. This new approach is a variant of SODA based on execution traces, which may come from any kind of SBSs. Conversely, SODA applies specifically on SCA SBSs using a set of scenarios and SCA techniques. In particular, in SOMAD, we specify a new set of metrics that apply to sequential association rules mined on execution traces whereas, in SODA, metrics apply to the concrete invocation of SBSs' interfaces using a set of scenarios. Figure 1 shows an overview of SOMAD. We emphasized in grey the two new steps specific to SOMAD and added to the SODA approach. The first new step (*Step 1. Metric Inference*) is supported by the creation of a set of hypotheses made from the textual description of SOA antipatterns. These hypotheses are used in this step to infer new metrics dedicated to the interpretation of sequential association rules. The second new step (*Step 4. Association Rule Mining*) aims to mine sequential association rules in execution traces of the targeted SBS. Sequential association rules are mined to discover interesting relations between services inside large amount of execution traces. In what follows, we first introduce key concepts of sequential association rules and then, present the overall process of SOMAD. Finally, we provide some implementation details.

A. Introduction to Sequential Association Rules

In the data mining field, Association Rule Mining (ARM) is a well-established method for discovering co-occurrences between attributes in the objects of a large data set [25]. Plain associations have the form $X \rightarrow Y$, where X and Y are sets of descriptors. Even though plain association rules could serve some relevant information, we are interested in preserving

the sequences of service invocations. We therefore adopt a variant called sequential association rules in which both X and Y become sequences of such descriptors (purchases by a customer, network alarms, or any other general kind of events). Moreover, in our case both sequences follow a temporal order whereas the left-hand side, aka the antecedent, comes before the right-hand one, the consequent. The rules that can be mined from traces bring to the foreground crucial information about the likelihood that services appear together in an execution trace and, more importantly, in a specific order. Thus, a possible rule could be : *ServiceA* , *ServiceB* implies *ServiceC* meaning that after executing A and then B, there are good chances to see C in the trace. For the sake of clarity, we limited the sizes of antecedent/consequent, nevertheless, each side of a rule can comprise larger sequences. The strength of the rule is measured by the *confidence* metric: In probabilistic terms, it measures the conditional probability of C appearing down the line. Beside that, the significance of a rule, i.e. how many times it appears in the data, is provided by its *support* measure. To ensure only rules of potentially high interestingness are mined, the mining task is tuned by minimal thresholds to output only the sufficiently high scores for both metrics.

B. SOMAD Process

Step 1. Metrics Inference: A set of metrics dedicated to the interpretation of sequential association rules are inferred from a set of three hypotheses made using the textual description of SOA antipatterns (Table I).

These hypotheses correspond to heuristics that allow the interpretation of sequential association rules in order to identify properties relevant to SOA antipatterns. After a careful examination of antipattern textual descriptions, we observe that SOA antipatterns can be specified in terms of coupling¹ and cohesion².

Hypothesis 1. *If a service A implies a service B with a high support and a high confidence, then A and B are tightly coupled.*

Hypothesis 2. *If a service appears in the consequent (antecedent) parts for a high number of associations, then it has high incoming (outgoing) coupling.*

The above hypotheses qualify the coupling between two

¹The coupling basically refers to the degree a services relies on others [26]

²Recall cohesion measures the relatedness between responsibilities of a service [26].

Multi-Service, *a.k.a* God Object corresponds to a service that implements a **multitude of methods** related to different business and technical abstractions. This aggregate too much into a single service, such a service is not easily reusable because of the **low cohesion** of its methods and is often unavailable to end-users because of its overload, which may induce a high response time [27].

Tiny Service is a small service with **few methods**, which only implements part of an abstraction. Such service often requires **several coupled services** to be used together, resulting in higher development complexity and reduced usability. In the extreme case, a Tiny Service will be limited to **one method**, resulting in many services that implement an overall set of requirements [27].

Chatty Service corresponds to a set of services that exchange a **lot of small data** of primitive types. The Chatty Service is also characterized by a **high number of method invocations**. Chatty Service chats a lot with each other [27].

The Knot is a **set of very low cohesive** services, which are tightly coupled. These services are thus less reusable. Due to this complex architecture, the availability of these services can be low, and their response time high [28].

Bottleneck Service is a service that is **highly used** by other services or clients. It has a **high incoming and outgoing coupling**. Its response time can be higher because it may be used by too many external clients, for which clients may need to wait to get access to the service. Moreover, its availability may also be low due to the traffic.

Service Chain, *a.k.a*. Message Chain in OO systems, corresponds to a **chain of services**. The Service Chain appears when clients request **consecutive service invocations** to fulfill their goals. This kind of **dependency chain** reflects the action of invocation in a transitive manner.

Table I: List of SOA Antipatterns [2]

specific services and overall incoming/outgoing coupling. The cohesion is also widely used in SOA antipattern descriptions.

Hypothesis 3. *If the number of different methods of a given service is similar to the number of different partners (Hypothesis 2) it has in the service rules, then the service is not cohesive.*

Based on the above three hypotheses, we have created domain specific metrics to help us explore the antipattern manifestations that are hidden in the sequential association rules. We use the DSL we defined in [2] to combine them. Metrics are presented in Figure 2. In the figure, standard mathematical notations are used whenever possible and extended if necessary. Thus, association rules are visualized by $(X \rightarrow Y)$ with X and Y represent the antecedent and the consequent parts, respectively. K, L are partner services. AR stands for the overall set of association rules while AR_s and AR_m being subsets targeting association rules at service / method level, respectively. M_S denotes the methods of a given service S . Finally, we use non-standard symbols for sequence operations: \square is the sequence constructor, \cup stand for append on sequences; \subseteq denotes the sub-sequence-of relationship; and $A \leq B$ means the service/method A appears inside B. Metrics can be combined to define other metrics.

Step 2. Specification of SOA antipatterns: The combination of metrics defined in the previous step allows the specification of SOA antipatterns in the form of sets of rules, called *rule cards*.

For the individual metrics and combinations thereof, the values that trigger the detailed examination of a case are not fixed be-

```

1 RULE_CARD: MultiService {
2   RULE: MultiService{INTER LowCohesion ManyMethods ManyMatches};
3   RULE: LowCohesion{COH LOW};
4   RULE: ManyMethods{NM HIGH};
5   RULE: ManyMatches{NMA HIGH};
6 };
      (a) Multi Service

1 RULE_CARD: TinyService {
2   RULE: TinyService{INTER HighOutgoingCoupling FewMethods};
3   RULE: HighOutgoingCoupling{OC HIGH};
4   RULE: FewMethods{NM LOW};
5 };
      (b) Tiny Service

1 RULE_CARD: ChattyService {
2   RULE: ChattyService{INTER ManyPartners ManyMatches};
3   RULE: ManyPartners{NDP VERY HIGH};
4   RULE: ManyMatches{NMA VERY HIGH};
5 };
      (c) Chatty Service

1 RULE_CARD: BottleNeck {
2   RULE: BottleNeck{INTER HighOutgoingCoupling HighIncomingCoupling};
3   RULE: HighOutgoingCoupling{OC HIGH};
4   RULE: HighIncomingCoupling{IC HIGH};
5 };
      (d) BottleNeck Service

1 RULE_CARD: KnotService {
2   RULE: KnotService{INTER LowCohesion HighCrossInvocation};
3   RULE: LowCohesion{COH LOW};
4   RULE: HighCrossInvocation{CID HIGH};
5 };
      (e) Knot Service

1 RULE_CARD: ServiceChain {
2   RULE: ServiceChain{HighTransitiveCoupling};
3   RULE: HighTransitiveCoupling{TC HIGH};
4 };
      (f) Service Chain

```

Figure 3: Rule Cards

forehand. Instead, we use a boxplot-based statistical technique that exploits the distribution of all values across the sets of services, methods, and rules. Moreover, the computed values are further weighted using the quality metrics for associations, i.e. support and confidence, so that the strongest rules could be favored.

The *rule cards* used to specify SOA antipatterns are presented in Figure 3. As an example, the rule card corresponding to the Tiny Service specification (Figure 3(b)) is composed of three rules. The first one (line 2) is the intersection of two rules (lines 3, 4), which define two metrics: a high Outgoing Coupling (OC) and a low Number of Method (NM).

Step 3. Generation of detection algorithms: This step stays unchanged from SODA, as described in Section II-C.

Step 4. Association Rule Mining: Execution traces are analyzed to extract the sequential association rules.

Association rules are extracted from a collection of sequence-shaped transactions with respect to a minimal support and a minimal confidence threshold. A transaction is a time-ordered set of different services and method calls. Recall that the support of a pattern, i.e. sequence of items (services or service methods), reflects the overall percentage of transactions that contain the pattern, whereas the confidence measures the likelihood of the consequent following the occurrence of the antecedent in a transaction. For our experiments (see next section) we set the values of the thresholds to 40% and 60%, respectively. The choice of these values does not follow any specific indication, general law from ARM or deeper insight into the SBS architecture. As our approach is at its exploratory stage, we were only guided by the need to filter

<p>Number of Matches (NMA(S)) : $\#\{X \rightarrow Y \in AR_s \mid S \triangleleft (X \uplus Y)\}$</p> <p>Follows the number of rules where a service appears, either on the left- or on the right-hand side.</p>
<p>Number of Diff. Partners (NDP(S)) : $\#\{K \mid X \rightarrow Y \in AR_s, S \triangleleft X, K \triangleleft Y\} + \#\{K \mid X \rightarrow Y \in AR_s, S \triangleleft Y, K \triangleleft X\}$</p> <p>Indicates how many different partners a service has. Spelled differently, the metric determines whether the service communicates intensively with surrounding services or not.</p>
<p>Number of Methods (NM(S)) : $\#\{K \mid X \rightarrow Y \in AR_m, K \in M_s, K \triangleleft (X \uplus Y)\}$</p> <p>Counts the number of occurrences of the methods from a service. The counting for this metric focuses on method rules.</p>
<p>Cohesion (COH(S)) : $\frac{NDP(S)}{NM(S)}$</p> <p>Assesses the ration between the numbers of partner services and of the available methods, respectively.</p>
<p>Cross Invocation Dependencies (CID(S_a, S_b)) : $\#\{X \rightarrow Y \in AR_s \mid S_a \triangleleft X, S_b \triangleleft Y\} + \#\{X \rightarrow Y \in AR_s \mid S_a \triangleleft Y, S_b \triangleleft X\}$</p> <p>Is paramount for the detection of antipattern manifestations dug into the traces. Indeed, the metric would explore the typical interactions between services while ignoring less frequent ones (absent from the mining method output due to the support threshold). To retrieve this information CID counts all association rules where a service A (S_a) is present in the antecedent and a service B (S_b) in the consequent or vice versa.</p>
<p>Incoming Coupling (IC(S)) : $\sum_{L \in \{K \mid X \rightarrow Y \in AR_s, K \triangleleft X, S \triangleleft Y\}} \frac{CID(S, L)}{NDP(S)}$</p> <p>Count how many times a service is used. Yet instead of merely counting a unit for each partner service, we use a contextual value: $\frac{CID(S, X)}{NDP(S)}$ where X is the partner service. Thus, the larger the portion of the partner service in the overall number of partners of S, the higher the coupling.</p>
<p>Outgoing Coupling (OC(S)) : $\sum_{L \in \{K \mid X \rightarrow Y \in AR_s, S \triangleleft X, K \triangleleft Y\}} \frac{CID(L, S)}{NDP(S)}$</p> <p>The same principle as for IC applies, yet in a dual manner. count how many times the argument service uses other services.</p>
<p>Transitive Coupling (TC(S_a, S_b)) : $\#\{K \mid X \rightarrow Y \in AR_s, S_a \triangleleft X, S_b \triangleleft Y, ([S_a, K] \subseteq X \vee [K, S_b] \subseteq Y)\}$</p> <p>Is aimed at spotting a common antipattern in SOA, the <i>Service Chain</i> (see above). First, observe that the fact that a pair of services does not communicate directly does not mean that they are not coupled: that is the founding idea of <i>Service Chain</i> and of our metric too. To identify transitive coupling manifestations within the extracted association rules, which is more challenging than the direct coupling, we need to capture describing a chain: e.g. a service S_a is in the antecedent of a rule, S_b is in the consequent of another rule and both rules are connected by means of a third service K that appears in the consequent of the first rule and in the antecedent of the second one. Longer chains are possible as well. Thus, in the basic case, we could have [A]→[B] and [B]→[C], with this configuration, A and C are not directly coupled but if C fails, there are good chances that A and B would fail too.</p>

Figure 2: Metrics \uplus : append on sequences; \subseteq : sub-sequence-of relationship; and $A \triangleleft B$: A appears inside B.

out all spurious associations while still keeping enough rules to represent the most significant calls (regulated via the support threshold). Moreover, we needed enough confidence in the threshold to make appear the most significant alternatives (rule consequent) for the termination of a specific sequence of calls (rule antecedent) while suppressing the less significant ones. Thus, we made several tries with different values for both thresholds and observed the size of the result. The current values seem to offer the best trade-off between size and completeness of scenarios.

To extract the association rules from execution traces, two choices were possible. On one hand, sequential pattern mining and rule mining algorithms have been designed for structures that are slightly more general than the ones used here. In fact, sequential patterns are defined on transactions that represent sequences of sets. Efficient sequential pattern miners have been published, e.g. the PrefixSpan method [29]. On the other hand, execution traces do not compile to fully-blown sequential transactions as the underlying structures are mere sequences of individual elements. Such data has been known

since at least the mid-90s but received less attention by the data mining community, arguably because it is less challenging to mine. In contrast, many practical applications have been reported where such data arise, inclusive software log mining (see Section II). In the general data mining literature, mining from pure sequences, as opposed to sequences made of sets, has been addressed under the name of episode mining [30]. Episodes are made of *events* and in a sense, service calls are events. Arguably the largest body of knowledge on the subject belongs to the web usage mining field: The input data is again a system trace, yet this time the trace of requests sent to a web server [31]. It is noteworthy that sequential patterns are more general than the pure sequence ones, hence mining algorithms designed for the former might prove to be less efficient when applied to the latter (as additional steps might be required for listing all significant set). Nevertheless, to jump-start our experimental study, we used a sequential pattern/rule miner that has the advantage to be freely available on the web³. Although it has not been optimized for pure sequences its

³<http://www.philippe-fournier-viger.com/spmf/>

performances are more than satisfactory. However, in a follow up study, we shall be implementing a more advanced method specifically targeting traces (see Section V). In summary, at the end of this, we have extracted the relevant relationship between services in the form of sequential association rules.

Step 5. Detection of SOA antipatterns

The last step of SOMAD consists in applying the detection algorithms generated in Step 3 on the sequential association rules mined in Step 4. At the end of this step, services in the SBS suspected to be involved in an antipattern are identified.

C. Implementation details

In this subsection, we present implementation details for other steps that may support the SOMAD approach.

Generation of Execution Traces. In case execution traces are not available, this step allows their generation.

If the targeted SBS does not produce qualitative execution traces that contain all the required information, we have to instrument it to do so. Such traces enable low-tech application debugging support whenever debuggers are unavailable or inapplicable (frequently the case with SOA environments). Trace producing can introduce source code obfuscation, however, it may have some benefits in terms of design quality as the code must be well mastered in order to correctly instrument. This technique of tracing is the most common, however if the source code is not available another technique consists in instrumenting the running environment of the SBS. This environment can be the virtual machine, the web server, or the operating system. For example, LTTng [32] instruments Linux to produce traces with a very low overhead.

To ease automated processing of traces, we provide a template (see Figure 4) that is a good trade-off between simplicity and information content. In this template, a method invocation generates two lines, an opening and a closing one with belonging customer identification (IP address) and a timestamp.

SBSs often contain built-in tracing systems, however, these built-in mechanisms can be very distant from our template. Thus, we made SOMAD adaptable by a simple-regular expression based-domain specific language (Figure 4). Therefore, traces can be ordered differently.

```
IP timestamp void methodA.ServiceA();
IP timestamp void methodB.ServiceB();
IP timestamp end void methodB.ServiceB();
IP timestamp end void methodA.ServiceA();

time{^w+\\s\\d\\s\\d\\d\\d\\d\\d+} || end{end}
method{^[^]*.*$} || service{^[^]*.*$}
customer{\\b\\d{1,3}\\d{1,3}\\d{1,3}\\d{1,3}\\b}
line{ *customer *time *(end)? *method.service *}
```

Figure 4: Trace template and associated DSL definition

Collecting and Aggregating traces. The goal here is to download all distributed trace files and aggregate them in a single one.

Traces are generated by several services composing the SBS

and then collected and aggregated. This is a very important and non-trivial task [33]; indeed the dynamic and distributed nature of SBSs implies, at least, two challenges. The first one is related to the distributivity of SBSs; and therefore, of execution traces. Indeed, each service will generate its own execution traces in its own running environment. Therefore, we need to know the names and running place of each service and a mechanism to download / retrieve execution traces on each running environment. Moreover, services can be consumed by several customers at the same time, execution traces can be interleaved. To solve these problems we propose an approach inspired by A. Yousefi and K. Startipi [34]. We first gather all executions log files in one file. Then, we sort execution traces using their timestamp and exploit the caller-callee relationships determined by service and method names to identify blocks of concurrent traces.

Focus shift. This step explains why SOMAD performs better than SODA to identify truly harmful SOA antipatterns.

SOMAD hypothesises shift the focus of the antipattern search from pure architectural consideration to usage, thus neglecting the exact values of some basic metrics. It is a natural choice since SOMAD is an approach, which does not have access to exact values throughout service interfaces or implementation. Moreover, analyzing a system through its utilization will, as proven in the experimentation part, have better precision. Considering a service named *Half-Deprecated Service* composed of four methods: A, B, C and D. The methods C and D are outdated and the service still exposes them to ensure retro-compatibility. One method to compute the cohesion of a service is to count how many methods of this service are used during the same session by a unique user. Since half of the methods are outdated; is more likely that user will consume only two methods on the four available. Therefore, if the cohesion is computed based on the service's interface the result will be 0.5 (2/4) and this service will be considered as a suspicious service for SOA antipatterns relying on low cohesion. On the contrary, if the cohesion is computed using the execution traces based method the result will most likely be 1.0. Indeed, the unforeseen call to the deprecated methods will be deleted because of their low support or confidence. Therefore, SOMAD performs better than SODA to detect harmful and true positive SOA antipatterns.

IV. EXPERIMENTS

As a validation study, we apply SOMAD on two independently developed SBSs, *Home Automation* and *FraSCAti* [8]. *Home Automation* is an SBS made of 13 services and selected for comparison with the outcome produced by SODA, the so far unique state-of-the-art tool for antipattern detection. Both tools were evaluated in terms of precision and recall, on one hand, and efficiency, on the other hand. We also apply SOMAD to *FraSCAti*, an SBS almost 10 times larger than *Home Automation*, which contains 91 components and 130 services.

A. Subjects

We apply SOMAD to detect six different SOA antipatterns described in Table I. In the description of each antipattern, we highlight in bold the characteristics relevant for their detection using our metrics.

B. Objects

A first round of experiments was performed on *Home Automation*, the same system used in the validation of SODA. *Home Automation* is an independently developed SCA-based system for remotely controlling basic household functions (i.e., temperature, electrical instruments, medical emergency support, etc.) in home care support for elderly. It also includes a set of 7 predefined scenarios for test and demonstration purposes. Two different versions of the system were used: the original version, made of 13 services, and an intentionally degraded version in which services have been modified and new ones added in order to inject some antipatterns. The changes were performed by a third-party to avoid bias in the results.

Given the lack of freely available SBSs, the second round was performed on FraSCAti [8], the runtime support of *Home Automation*. FraSCAti is also an SCA-based system made of more than 90 components and over 130 services scattered between components. A component exposes at least one service and services expose methods. Unlike *Home Automation*, FraSCAti does not have predefined scenarios—in reality it provides some unit tests, but not complete feature coverage. The detection was performed by instrumenting FraSCAti to produce execution traces as described in Section III-C. As FraSCAti is a runtime support for SOA systems, we loaded and ran diverse SBSs of different technologies (SCA, REST, Web Services, RMI-based) and then, use these systems to have a maximum feature coverage. The detection of SOA antipatterns in FraSCAti has been performed at the component level instead of service-level since the system architecture is documented at that level while the subsequent validation will be based on this documentation. Moreover, it was empirically established that SCA-based systems suffer from the same architectural flaws as pure SOA systems. Details on the systems including all the scenarios and involved services are available online at <http://sofa.uqam.ca/somad>.

C. Process

We applied SOMAD for the detection of the six SOA antipatterns on the two targeted SBSs. First, we run the seven scenarios of *Home Automation* on its two versions, and then the six scenarios of FraSCAti. Then, we recreated transactions from the execution traces and run our algorithm for rule generation, with a support of 40% and a confidence of 60%, the corresponding sequential association rules. The step that follows consisted in interpreting the generated association rules. For this purpose, we computed the metrics associated to conjectures that fit the textual descriptions of the six SOA antipatterns. After this step of interpretation, we obtained for each SBS the list of suspicious services involved in the antipatterns. Finally, we validated the detection results in terms of precision and recall by analyzing the suspicious services manually. Precision estimates the ratio of true antipatterns

identified among the detected antipatterns, while recall estimates the ratio of detected antipatterns among the existing antipatterns. This validation has been performed manually by an independent software engineer, whom we provided the descriptions of antipatterns, the two versions of the analyzed system *Home Automation*, and the system FraSCAti with a printed description of its architecture available online on the FraSCAti web site (<http://frascati.ow2.org>). For both systems, we compared the results with the ones obtained by SODA. For FraSCAti, we reported the detection results to their development team and got their feedback as a objective validation.

D. Results

Table II presents the results for the detection of the six SOA antipatterns on the original and evolved version of *Home Automation*. For each SOA antipattern, the table reports the version analysed of *Home Automation*, the services detected automatically by SOMAD, the services identified manually by the software engineer, the metric values, the recall and precision, the computation time, and finally, the F-measure [23], which corresponds to the harmonic mean of precision and recall. Similarly, Table III provides the detection results on FraSCAti. We recall that the metric values reported in the tables do not represent absolute values (e.g. for NM, the exact number of methods exposed), but rather elicit what we called the *usage representation* of a SBS. And in particular, the metric values are weighted by the fraction $\frac{\text{support}}{\text{confidence}}$ for highlighting most confident and supported association rules. Thus, a number of methods (NM) of 2 means that among the generated association rules, there are 2 methods that appear in the rules with a high support and confidence.

E. Details of the results

We present the detection results of SOMAD while comparing them to SODA, both on the system *Home Automation*. The results with SOMAD are quite similar to the ones obtained with SODA, except for *The Knot* and *Bottleneck Service* antipatterns.

For example, IMediator has been detected and identified as a *Multi Service*, both in SOMAD and SODA, because of its high number of methods ($NM \geq 2$), its high number of matches ($NMA \geq 3.8$) and its low cohesion ($COH \leq 0.5$). These metric values have been evaluated as high and low in comparison with the metric values of *Home Automation*. For example, for the metric NM, the boxplot estimates the high value of NM in *Home Automation* as equal to 2. Similarly, the detected *Tiny Service* has a very low number of methods ($NM \leq 1$) and a high outgoing coupling ($OC \geq 4$) according to the boxplot. In the original version of *Home Automation*, we did not detect any *Tiny Service*. An independent engineer extracted one method from IMediator and moved it into a new service named MediatorDelegate; this newly injected service has been detected as a *Tiny Service*. Two occurrences of *Chatty Service* have been discovered in *Home Automation*, both in SOMAD and SODA. PatientDAO and IMediator have a high number of matches ($NMA \geq 3.8$), which mean that the service *talks* too much, and they have a high number of different partners ($NDP \geq 0.6$).

PatientDAO has been detected as a *Knot* because it has

Antipattern Name		Automatically detected services	Manually identified services	SOMAD Metrics	Recall	Precision	Time	F ₁
Tiny Service	SODA	Mediator-Delegate	Mediator-Delegate	OC ≥ 4	[1/1] 100%	[1/1] 100%	0.194s	100%
	SOMAD	Mediator-Delegate		NM ≤ 1	[1/1] 100%	[1/1] 100%	0.077s	100%
Multi Service	SODA	IMediator	IMediator	NM ≥ 2 NMA ≥ 3.8	[1/1] 100%	[1/1] 100%	0.462s	100%
	SOMAD	IMediator		COH ≤ 0.5	[1/1] 100%	[1/1] 100%	0.050s	100%
Chatty Service	SODA	PatientDAO IMediator	PatientDAO	NMA ≥ 3.8	[2/2] 100%	[2/2] 100%	0.383s	100%
	SOMAD	PatientDAO IMediator	IMediator	NDP ≥ 0.6	[2/2] 100%	[2/2] 100%	0.077s	100%
The Knot	SODA	PatientDAO IMediator	PatientDAO	CID ≥ 2	[1/1] 100%	[1/2] 50%	0.412s	66.6%
	SOMAD	PatientDAO		COH ≤ 0.5	[1/1] 100%	[1/1] 100%	0.077s	100%
BottleNeck	SODA	IMediator PatientDAO	IMediator	IC ≥ 4	[2/2] 100%	[2/2] 100%	0.246s	100%
	SOMAD	IMediator PatientDAO SunSpotService	PatientDAO	OC ≥ 3	[2/2] 100%	[2/3] 66%	0.076s	79.5%
Chain Service	SODA	{IMediator, PatientDAO, SunSpotService, PatientDAO2}	{IMediator, PatientDAO, PatientDAO2}	LC ≥ 4	[3/3] 100%	[3/4] 75%	0.229s	85.7%
	SOMAD	{IMediator, PatientDAO, SunSpotService, PatientDAO2}			[3/3] 100%	[3/4] 75%	0.056s	85.7%
Average	SODA				100%	87.5%	0.231s	92.0%
	SOMAD				100%	90.1%	0.068s	94.2%

Table II: Results comparison between SODA & SOMAD on HomeAutomation

Antipattern Name		Automatically detected services	Manually identified services	SOMAD Metrics	Recall	Precision	Time	F ₁
Tiny Service	SODA	SCA-Parser	SCA-Parser	OC ≥ 3	[1/1] 100%	[1/1] 100%	0.083s	100%
	SOMAD	SCA-Parser		NM ≤1	[1/1] 100%	[1/1] 100%	0.066	100%
Multi Service	SODA	juiliac Explorer-GUI	Explorer-GUI	NDP ≥ 24 NMA ≥ 70	[1/1] 100%	[1/2] 50%	0.462s	66.67%
	SOMAD	Explorer-GUI		COH ≤ 0.5	[1/1] 100%	[1/1] 100%	0.050s	100%
Chatty Service	SODA	<i>not present</i>	<i>not present</i>	NMA ≥ 70	[0/0] 100%	[0/0] 100%	0.97s	100%
	SOMAD	<i>not present</i>		NDP ≥24	[0/0] 100%	[0/0] 100%	0.77s	100%
The Knot	SODA	SCA-Parser SCA-Composite	SCA-Parser	CID ≥ 25	[1/1] 100%	[1/2] 50%	1.041s	66.6%
	SOMAD	SCA-Parser		COH ≤ 0.2	[1/1] 100%	[1/1] 100%	0.7s	100%
BottleNeck	SODA	SCA-Composite SCA-Parser	SCA-Parser	IC ≥ 3	[2/2] 100%	[2/2] 100%	0.246s	100%
	SOMAD	SCA-Parser SCA-Composite Metamodel-Provider	SCA-Composite	OC ≥ 3	[2/2] 100%	[2/3] 66.67%	0.076s	80%
Chain Service	SODA	SCA-Parser Composite-Manager Processor	Composite Parser Composite-Manager	LC ≥ 5	[2/2] 100%	[2/3] 66.67%	0.758	80%
	SOMAD	Composite-Parser Composite-Manager			[2/2] 100%	[2/2] 100%	0.056s	100%
Average	SODA				100%	77.77%	0.707s	85.55%
	SOMAD				100%	94.44%	0.28s	96.6%

Table III: Results comparison between SODA & SOMAD on FraSCaTi

a high cyclic invocation dependencies ($CID \geq 2$) and a low cohesion ($COH \leq 0.5$). The metric CID allows the identification of cyclic invocation dependency. In *Home Automation*, the set of services PatientDAO1, PatientDAO2, PatientDAO3, PatientDAO4 are tightly coupled because each of them represents a part of a patient's information (name, address, phone number). Therefore, cyclic invocations between these services appear when information about a patient are requested. SOMAD does not report the false positive, IMediator, reported by SODA, and thus obtains a better precision for this antipattern.

Three services have detected as *BottleNeck Services*: IMediator PatientDAO, and SunSpotService because of their high outgoing and incoming coupling ($IC \geq 4$ and $OC \geq 3$). This time, it is SOMAD that reports the false positive, SunSpotService, and thus decreases its precision compared to SODA.

Finally, we detected both in SOMAD and SODA, the transitive chain of invocations IMediator \rightarrow PatientDAO \rightarrow PatientDAO2 \rightarrow SunSpotService ($LC \geq 4$). In both approaches, the false positive SunSpotService has been reported.

We now present the detection results of SOMAD on FraSCAti.

SCA-Parser is suspected to be a *Tiny Service* because it includes a low number of methods (NM equal 1) and a high outgoing coupling (OC equal 3). A manual code inspection of FraSCAti revealed that SCA-Parser contains only one interface method, named `parse(...)`. The development team of FraSCAti validated this antipattern. They indicated that this service can be invoked alone when only a reading of a SCA file is requested. However, FraSCAti performs more tasks that just reading an SCA file, and these other tasks are performed by other services such as AssemblyFactory. This explains the high outgoing coupling.

SOMAD did not detect any *Multi Service* in FraSCAti. However, the manual inspection of FraSCAti allowed the identification of the component Explorer-GUI as a *Multi Service*. The FraSCAti development team confirmed that this component uses a high number of services provided by other components of FraSCAti. Indeed, this component encapsulates the graphical interface of FraSCAti Explorer, which aims to provide an exhaustive interface of FraSCAti functionalities. SOMAD was not able to detect it because the execution scenarios did not involve the graphical interface of FraSCAti Explorer.

SOMAD did not detect any *Chatty Service* in FraSCAti. No service has a very high number of matches (NMA) and a very high number of different partners (NDP), respectively higher than 70 and 24. This means that no service appears more than 70 times in the set of association rules and communicates with more than 24 different other services. The manual code inspection confirmed also that there was no *Chatty Service* in FraSCAti.

The component Metamodel-provider is suspected to be part of a Knot because of its low cohesion ($COH \leq 0.2$) and its very high cyclic invocation dependencies ($CID \geq 25$).

The validation by the FraSCAti team has only confirmed that this component was implemented by many other components, but they did not agree on the specification of this antipattern. However, the independent software engineer validated this detection.

SOMAD detected three occurrences of the *BottleNeck Service* antipattern, SCA-Parser, Composite-Parser, and Metamodel-provider, the last of which was identified as a false positive. These services have been identified as *BottleNeck Services* because they have a high outgoing and incoming coupling (OC and $IC \geq 3$). The FraSCAti development team confirmed that SCA-Parser is highly used by other services.

Finally, Composite-Parser has been detected and identified as a *Chain Service*, whereas Composite-Manager is a false positive. Composite-Parser is involved in a long transitive chain of invocations ($LC \geq 4$). The FraSCAti development team validated this antipattern and indicated that this service uses a delegation chain to perform its behavior.

We can observe that Composite-Parser and SCA-Parser are very suspicious services. They are both involved in two antipatterns. These services are highly coupled with other services and in particular and are part of a long transitive invocation chain. The presence of such antipatterns in a system is not surprising because there is no other way to develop a parser without introducing a high coupling and high transitivity.

In conclusion, FraSCAti is performing reasonably well towards the antipattern detection. Few services have been detected as antipatterns compared to the high number of FraSCAti components/services. Mainly, SCA-Parser is on the critical path of all processing performed by FraSCAti.

All detailed results on *Home Automation* and FraSCAti are provided on our website : <http://www.sofa.uqam.ca/somad>

F. Threats to validity

The main threat to the validity of our results is directed at their *external validity*, i.e., the possibility to generalize the current results to other SBSs. Given the lack of freely available systems, we done our best to obtain realistic systems such as FraSCAti and we experimented with two versions of *Home Automation*. However, we plan to run these experiments on other SBSs in the future, with special focus on SBSs implementing other SOA technologies, such as REST and Web services. For *internal validity*, the detection results depend on our conjectures. Although we did not perform our experiments on a representative set of antipatterns as with SODA, we obtained comparable results in terms of precision and recall. The subjective nature of interpreting the association rules and validating antipatterns is a threat to the *construct validity*. We control this threat by specifying our conjectures based on a literature review on antipatterns and by involving in our study an independent engineer and the FraSCAti development team. Finally, we minimize the threats to *reliability validity* by automating the generation of association rules.

V. CONCLUSIONS AND FUTURE WORK

The detection of SOA antipatterns is a crucial activity if we have to ensure the architectural and overall quality

of SBS. In this paper, we present a new and innovative approach, SOMAD, for the detection of antipatterns. The approach relies on two complementary techniques, from two thriving fields in software engineering, mining system traces and software measurement, respectively, both put in a SOA environment. More precisely, SOMAD detects SOA antipatterns by first mining association rules from execution traces and then filtering the encoded knowledge about interconnections between services by means of domain-specific metrics. The usefulness of SOMAD was demonstrated by applying it to two independently developed SBS. The results of this new approach incremented from SODA, the only and therefore, state-of-the-art tool, to fill its limitations were compared to SODA ones; and the outcomes show that SOMAD is a relevant approach. Its precision is better by a fairly comfortable margin ranging from 2.6% to 16.67% while keeping the recall to 100% and being—at least—2.5 times faster. Moreover, SOMAD uses execution traces that may come from any SOA technologies while SODA focuses on SCA SBSs.

We envision the application of SOMAD in the context of a large data center whereby the goal would be to optimize the data center communications. We shall investigate alternative mining techniques to refine our approach with additional information, e.g. directly extracting architectural overviews with graph pattern mining [35] or, detecting recurring patterns of anomalous behavior with rare pattern mining [36]. Finally, combining explicit semantic representations, i.e. in OWL ontologies, with powerful mining methods for heterogeneous labeled graphs (see [37]) seems to be a particularly promising track for the extraction of complex structural and/or behavioral antipatterns from execution traces.

ACKNOWLEDGMENT

The authors thank Phillipe Merle and Lionel Seinturier for their help in understanding FraSCaTi. This work is partly supported by the Natural Sciences and Engineering Research Council of Canada.

REFERENCES

- [1] T. Erl, *Service-oriented architecture: concepts, technology, and design*. Pearson Education India, 2006.
- [2] N. Moha, F. Palma, M. Nayrolles *et al.*, “Specification and detection of soa antipatterns,” in *ICSOC*, 2012.
- [3] P. Wolfgang, *Design patterns for object-oriented software development*. Reading, Mass.: Addison-Wesley, 1994.
- [4] R. Marinescu, “Detection Strategies: Metrics-based Rules for Detecting Design Flaws,” in *IEEE ICSM*, 2004.
- [5] N. Moha, Y.-G. Guéhéneuc, L. Duchien, and A.-F. L. Meur, “Decor: A method for the specification and detection of code and design smells,” *IEEE TSE*, vol. 36, no. 1, pp. 20–36, Jan. 2010.
- [6] N. Tsantalis, T. Chaikalis, and A. Chatzigeorgiou, “JDeodorant: Identification and Removal of Type-Checking Bad Smells,” in *CSMR 2008*, Apr. 2008, pp. 329–331.
- [7] M. Nayrolles, F. Palma, and Moha, “Soda : A tool for automatic detection of soa antipatterns,” in *ICSOC - Tool Paper*, 2012.
- [8] L. Seinturier, P. Merle, R. Rouvoy, D. Romero, V. Schiavoni, and J.-B. Stefani, “A component-based middleware platform for reconfigurable service-oriented architectures,” *SPE*, vol. 42, no. 5, pp. 559–583, 2012.
- [9] M. Kessentini, S. Vaucher, and H. Sahraoui, “Deviance From Perfection is a Better Criterion Than Closeness To Evil When Identifying Risky Code,” in *Proceedings of the IEEE/ACM ASE*. ACM, 2010, pp. 113–122.
- [10] M. Lanza and R. Marinescu, *Object-Oriented Metrics in Practice*. Springer-Verlag, 2006.
- [11] W. Brown, R. Malveau, H. McCormick III, and T. Mowbray, *Anti Patterns: Refactoring Software, Architectures, and Projects in Crisis*. John Wiley and Sons, 1998.
- [12] M. J. Fowler, K. Beck, J. Brant, W. Opdyke, and D. Roberts, *Refactoring: Improving the Design of Existing Code*. Addison-Wesley, 1999.
- [13] D. L. Settas, G. Meditskos, I. G. Stamelos, and N. Bassiliades, “SPARSE: A symptom-based antipattern retrieval knowledge-based system using Semantic Web technologies,” *ESA*, vol. 38, no. 6, pp. 7633–7646, June 2011.
- [14] S. Wong, M. Aaron, J. Segall, K. Lynch, and S. Mancoridis, “Reverse Engineering Utility Functions Using Genetic Programming to Detect Anomalous Behavior in Software,” in *Proceedings of WCRE*. IEEE Computer Society, 2010, pp. 141–149.
- [15] T. Parsons and J. Murphy, “Detecting performance antipatterns in component based enterprise systems,” *JOT*, vol. 7, no. 3, pp. 55–90, April 2008.
- [16] P. Tonella and M. Ceccato, “Aspect mining through the formal concept analysis of execution traces,” in *Proceedings of the 11th WCRE 2004*. IEEE Computer Society, 2004, pp. 112–121.
- [17] A. Khan, A. Lodhi, V. Köppen, G. Kassem, and G. Saake, “Applying process mining in soa environments,” in *Proceedings of the 2009 ICSOC*. Springer-Verlag, 2009, pp. 293–302.
- [18] M. J. Asbagh and H. Abolhassani, “Web service usage mining: mining for executable sequences,” in *Proceedings of WSEAS ICACSE - Volume 7*, ser. ACS’07. Stevens Point, Wisconsin, USA: WSEAS, 2007, pp. 266–271.
- [19] S. Dustdar and R. Gombotz, “Discovering web service workflows using web services interaction mining,” *IJBPM*, vol. 1, pp. 256–266(11), 27 February 2007.
- [20] H. Safyallah and K. Sartipi, “Dynamic analysis of software systems using execution pattern mining,” *ICPC*, vol. 0, pp. 84–88, 2006.
- [21] A. Yousefi and K. Sartipi, “Identifying distributed features in soa by mining dynamic call trees,” in *ICSM*. IEEE, 2011, pp. 73–82.
- [22] B. Upadhyaya, R. Tang, and Y. Zou, “An approach for mining service composition patterns from execution logs,” *Journal of Software : Evolution and Process*, 2012.
- [23] J. K.-Y. Ng, Y.-G. Guéhéneuc, and G. Antoniol, “Identification of behavioural and creational design motifs through dynamic analysis,” *JSMRP*, vol. 22, no. 8, pp. 597–627, 2010.
- [24] L. Hu and K. Sartipi, “Dynamic analysis and design pattern detection in java programs,” in *Proceedings of SEKE’2008*, 2008, pp. 842–846.
- [25] G. Piatetsky-Shapiro, “Discovery, analysis, and presentation of strong rules,” *KDD*, pp. 229–238, 1991.
- [26] W. Stevens, G. Myers, and L. Constantine, “Structured design,” *IBM Systems Journal*, vol. 13, no. 2, pp. 115–139, 1974.
- [27] B. Dudley, S. Asbury, J. Krozak, and K. Wittkopf, *J2EE AntiPatterns*. John Wiley & Sons Inc, 2003.
- [28] A. Rotem-Gal-Oz, E. Bruno, and U. Dahan, *SOA Patterns*. Manning Publications Co., 2012.
- [29] J. Pei, J. Han, B. Mortazavi-Asl, Wang *et al.*, “Mining sequential patterns by pattern-growth: The prefixspan approach,” vol. 16, no. 11. IEEE, 2004, pp. 1424–1440.
- [30] H. Mannila, H. Toivonen, and A. Inkeri Verkamo, “Discovery of frequent episodes in event sequences,” *Data Mining and Knowledge Discovery*, vol. 1, pp. 259–289, 1997.
- [31] J. Pei, J. Han, B. Mortazavi-Asl, and H. Zhu, “Mining access patterns efficiently from web logs,” *Knowledge Discovery and Data Mining. Current Issues and New Applications*, pp. 396–407, 2000.
- [32] P.-m. Fournier and M. R. Dagenais, “Combined Tracing of the Kernel and Applications with LTTng,” in *Linux Symposium*, 2009.
- [33] N. Wilde, S. Simmons, M. Pressel, and J. Vandeville, “Understanding features in SOA,” in *Proceedings of SDSOA ’08*. New York, New York, USA: ACM Press, May 2008, p. 59.
- [34] A. Yousefi and K. Sartipi, “Identifying distributed features in SOA by mining dynamic call trees,” in *ICSM*, Sep. 2011, pp. 73–82.
- [35] D. Chakrabarti and C. Faloutsos, “Graph mining: Laws, generators, and algorithms,” *ACM Computing Surveys (CSUR)*, vol. 38, no. 1, p. 2, 2006.
- [36] L. Szathmary, A. Napoli, and P. Valtchev, “Towards rare itemset mining,” in *Proc. of the 19th IEEE Intl. ICTAI’07*. IEEE Computer Society, 2007, pp. 305–312.
- [37] M. Adda, P. Valtchev, R. Missaoui, and C. Djeraba, “A framework for mining meaningful usage patterns within a semantically enhanced web portal,” in *Proc. of C3S2E ’10*. ACM, 2010, pp. 138–147.